

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
import numpy as np
import pandas as pd

import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline

import plotly.graph_objects as go
import plotly.express as px

import seaborn as sns

## Models
import sklearn
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier

## Model evaluators
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.model_selection import RandomizedSearchCV, GridSearchCV
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import precision_score, recall_score, f1_score
from sklearn.metrics import RocCurveDisplay
```

```
df = pd.read_csv("/content/drive/MyDrive/heart_disease_project/heart_disease/heart-disease.csv")
df.shape
```

(303, 14)

```
df.head(10)
```

(303, 14)

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1
5	57	1	0	140	192	0	1	148	0	0.4	1	0	1	1
6	56	0	1	140	294	0	0	153	0	1.3	1	0	2	1
7	44	1	1	120	263	0	1	173	0	0.0	2	0	3	1
8	52	1	2	172	199	1	1	162	0	0.5	2	0	3	1
9	57	1	2	150	168	0	1	174	0	1.6	2	0	2	1

```
df.target.value_counts()
```

(303, 14)

target	count
1	165
0	138

```
data_to_plot = df['target'].value_counts().reset_index()
data_to_plot.columns = ['target', 'count']
```

```
fig = go.Figure(
    data=[
        go.Bar(
```

```

        x=data_to_plot['target'],
        y=data_to_plot['count'],
        marker_color='indigo'
    )
],
layout_title_text="target counts "
)


fig.update_layout(
    xaxis_title="targets",
    yaxis_title="counts",
    title_font_size=20,
    xaxis_tickangle=40,
    height=700,
    width=900
)

fig.show()

```



target counts



```
df.info()
```



```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column        Non-Null Count  Dtype  
---  --
 0   age           303 non-null   int64  
 1   sex           303 non-null   int64  
 2   cp            303 non-null   int64  
 3   trestbps      303 non-null   int64  
 4   chol          303 non-null   int64  
 5   fbs           303 non-null   int64  
 6   restecg       303 non-null   int64  
 7   thalach       303 non-null   int64  
 8   exang         303 non-null   int64  
 9   oldpeak       303 non-null   float64 
10   slope         303 non-null   int64  
11   ca            303 non-null   int64  
12   thal          303 non-null   int64  
13   target        303 non-null   int64  
dtypes: float64(1), int64(13)
memory usage: 33.3 KB

```

```
df.describe()
```

	age	sex	cp	trestbps	chol	fb	restecg	thalach	exang	oldpeak	slope
<b>count</b>	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
<b>mean</b>	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.528053	149.646865	0.326733	1.039604	1.399340
<b>std</b>	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.525860	22.905161	0.469794	1.161075	0.616226
<b>min</b>	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.000000	0.000000	0.000000
<b>25%</b>	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	133.500000	0.000000	0.000000	1.000000
<b>50%</b>	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	153.000000	0.000000	0.800000	1.000000
<b>75%</b>	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	1.000000	166.000000	1.000000	1.600000	2.000000
<b>max</b>	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.000000	6.200000	2.000000

### comapring features

```
df.sex.value_counts()
```

```

count
sex
1    207
0     96

dtype: int64

```

```
pd.crosstab(index=df.target, columns=df.sex)
```

```

sex  0    1
target
0    24  114
1    72   93

```

```
crosstab = pd.crosstab(index=df['target'], columns=df['sex'])
```

```
crosstab = crosstab.reset_index()
```

```
crosstab_melted = crosstab.melt(id_vars='target', var_name='sex', value_name='count')
```

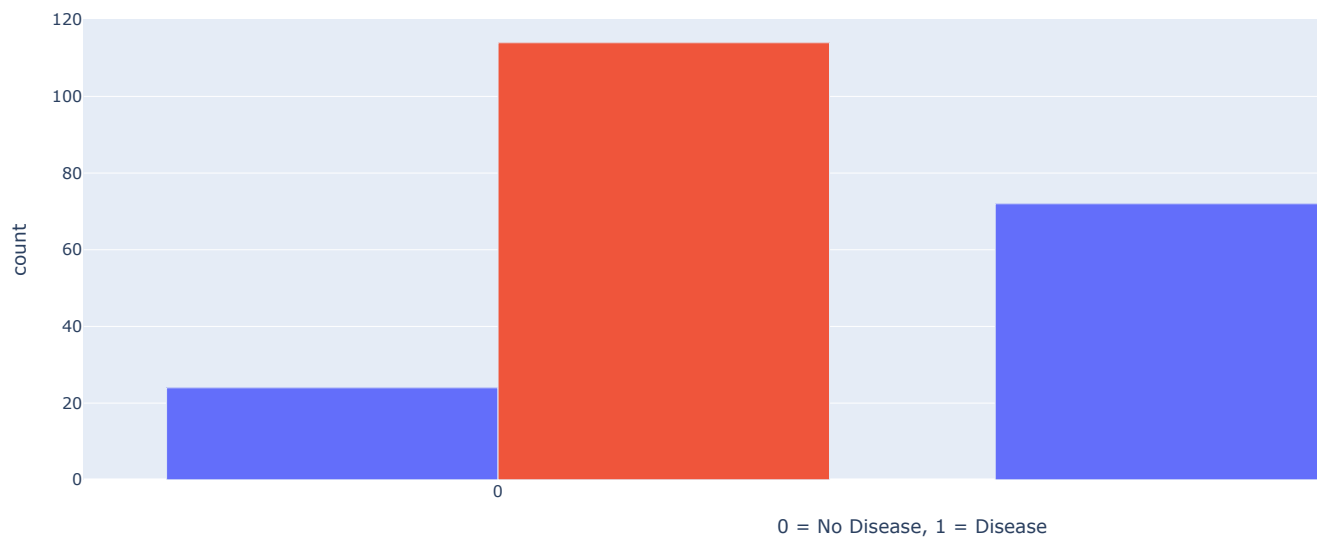
```

fig = px.bar(crosstab_melted, x='target', y='count', color='sex', barmode='group', title = "Heart Disease Frequency as per Sex")
fig.update_layout(
    xaxis_title="0 = No Disease, 1 = Disease",
    yaxis_title='count'
)
fig.show()

```




Heart Disease Frequency as per Sex



```
fig = px.scatter(  
    df,  
    x='age',  
    y='thalach',  
    color='target',  
    color_discrete_map={'Disease': 'salmon', 'No Disease': 'lightblue'},  
    labels={  
        'age': 'Age',  
        'thalach': 'Max Heart Rate',  
        'target': 'Condition'  
    },  
    title='Heart Disease in Function of Age and Max Heart Rate'  
)  
fig.show()
```



Heart Disease in Function of Age and Max Heart Rate



```
df.head()
```



	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

```

cp_target_ct = pd.crosstab(df['cp'], df['target'])

cp_target_ct = cp_target_ct.reset_index()

cp_target_ct.columns = ['Chest Pain Type', 'No Disease', 'Disease']

cp_target_melted = cp_target_ct.melt(id_vars='Chest Pain Type',
                                     value_vars=['No Disease', 'Disease'],
                                     var_name='Heart Disease',
                                     value_name='Frequency')

fig = px.bar(cp_target_melted,
             x='Chest Pain Type',
             y='Frequency',
             color='Heart Disease',
             barmode='group',
             title='Heart Disease Frequency Per Chest Pain Type',
             color_discrete_map={'No Disease': 'lightblue', 'Disease': 'salmon'})

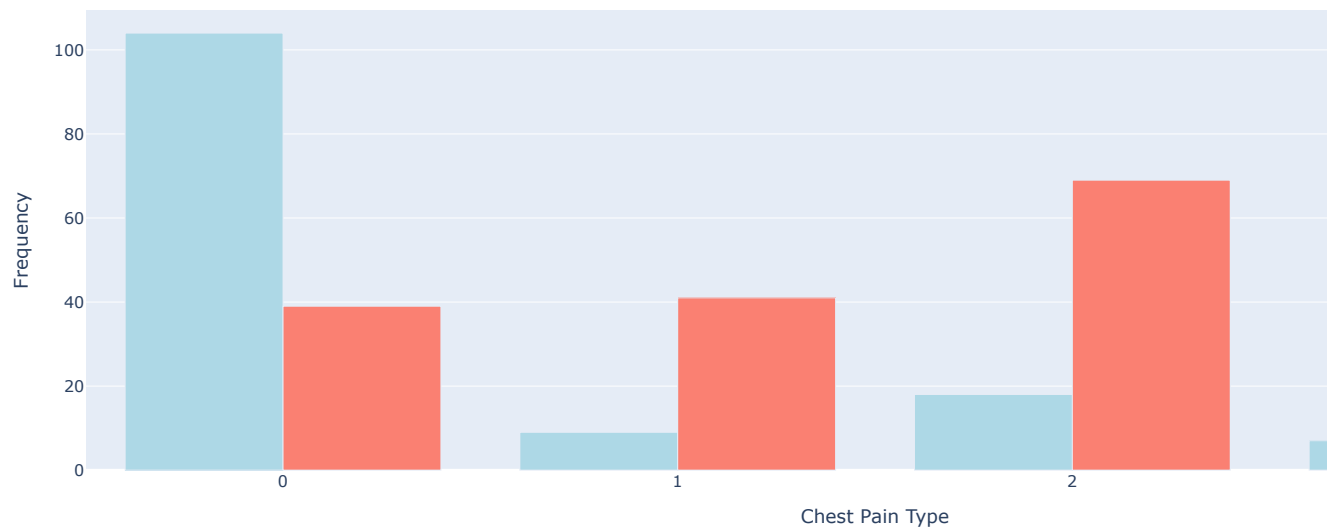
fig.update_layout(xaxis_title='Chest Pain Type',
                  yaxis_title='Frequency')

fig.show()

```




Heart Disease Frequency Per Chest Pain Type



```

corr_matrix = df.corr()
corr_matrix

```



	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	
<b>age</b>	1.000000	-0.098447	-0.068653	0.279351	0.213678	0.121308	-0.116211	-0.398522	0.096801	0.210013	-0.168814	0.276326	C
<b>sex</b>	-0.098447	1.000000	-0.049353	-0.056769	-0.197912	0.045032	-0.058196	-0.044020	0.141664	0.096093	-0.030711	0.118261	C
<b>cp</b>	-0.068653	-0.049353	1.000000	0.047608	-0.076904	0.094444	0.044421	0.295762	-0.394280	-0.149230	0.119717	-0.181053	-C
<b>trestbps</b>	0.279351	-0.056769	0.047608	1.000000	0.123174	0.177531	-0.114103	-0.046698	0.067616	0.193216	-0.121475	0.101389	C
<b>chol</b>	0.213678	-0.197912	-0.076904	0.123174	1.000000	0.013294	-0.151040	-0.009940	0.067023	0.053952	-0.004038	0.070511	C
<b>fbs</b>	0.121308	0.045032	0.094444	0.177531	0.013294	1.000000	-0.084189	-0.008567	0.025665	0.005747	-0.059894	0.137979	-C
<b>restecg</b>	-0.116211	-0.058196	0.044421	-0.114103	-0.151040	-0.084189	1.000000	0.044123	-0.070733	-0.058770	0.093045	-0.072042	-C
<b>thalach</b>	-0.398522	-0.044020	0.295762	-0.046698	-0.009940	-0.008567	0.044123	1.000000	-0.378812	-0.344187	0.386784	-0.213177	-C
<b>exang</b>	0.096801	0.141664	-0.394280	0.067616	0.067023	0.025665	-0.070733	-0.378812	1.000000	0.288223	-0.257748	0.115739	C
<b>oldpeak</b>	0.210013	0.096093	-0.149230	0.193216	0.053952	0.005747	-0.058770	-0.344187	0.288223	1.000000	-0.577537	0.222682	C
<b>slope</b>	-0.168814	-0.030711	0.119717	-0.121475	-0.004038	-0.059894	0.093045	0.386784	-0.257748	-0.577537	1.000000	-0.080155	-C
<b>ca</b>	0.276326	0.118261	-0.181053	0.101389	0.070511	0.137979	-0.072042	-0.213177	0.115739	0.222682	-0.080155	1.000000	C
<b>thal</b>	0.068001	0.210041	-0.161736	0.062210	0.098803	-0.032019	-0.011981	-0.096439	0.206754	0.210244	-0.104764	0.151832	1
<b>target</b>	-0.225439	-0.280937	0.433798	-0.144931	-0.085239	-0.028046	0.137230	0.421741	-0.436757	-0.430696	0.345877	-0.391724	-C

```
corr_matrix = df.corr()

fig = go.Figure(data=go.Heatmap(
    z=corr_matrix.values,
    x=corr_matrix.columns,
    y=corr_matrix.index,
    colorscale='YlGnBu',
    zmin=-1, zmax=1,
    text=np.round(corr_matrix.values, 2),
    texttemplate="%{text}",
    colorbar=dict(title="Correlation")
))

fig.update_layout(
    title="Correlation Matrix of Features",
    xaxis_title="Features",
    yaxis_title="Features",
    width=900,
    height=700,
)

fig.show()
```



```
np.random.seed(42)

X_train, X_test, y_train, y_test = train_test_split(X,
                                                    y,
                                                    test_size = 0.2)
```

```
X_train.head()
```

```
↗
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
<b>132</b>	42	1	1	120	295	0	1	162	0	0.0	2	0	2
<b>202</b>	58	1	0	150	270	0	0	111	1	0.8	2	0	3
<b>196</b>	46	1	2	150	231	0	1	147	0	3.6	1	0	2
<b>75</b>	55	0	1	135	250	0	0	161	0	1.4	1	0	2
<b>176</b>	60	1	0	117	230	1	1	160	1	1.4	2	2	3

```
X_test.head()
```

```
↗
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
<b>179</b>	57	1	0	150	276	0	0	112	1	0.6	1	1	1
<b>228</b>	59	1	3	170	288	0	0	159	0	0.2	1	0	3
<b>111</b>	57	1	2	150	126	1	1	173	0	0.2	2	1	3
<b>246</b>	56	0	0	134	409	0	0	150	1	1.9	1	2	3
<b>60</b>	71	0	2	110	265	1	0	130	0	0.0	2	1	2

```
y_train, len(y_train)
```

```
↗ (array([1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1,
          1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 0,
          1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1,
          0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0,
          0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0,
          1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1,
          1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1,
          1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0,
          0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 1,
          1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1,
          1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1,
          1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1]),
        242)
```

```
y_test, len(y_test)
```

```
↗ (array([0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0,
          0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1,
          1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0]),
        61)
```

```
models = {"KNN": KNeighborsClassifier(),
          "Logistic_Regression": LogisticRegression(max_iter=50,solver='liblinear'),
          "Random_Forest": RandomForestClassifier()}
```

```
def fit_and_score(models, X_train, X_test, y_train, y_test):
```

```
    np.random.seed(42)
```

```
    model_scores = {}
```

```
    for name, model in models.items():
```

```
        model.fit(X_train, y_train)
```

```
        model_scores[name] = model.score(X_test, y_test)
```

```
    return model_scores
```



```

model_scores = fit_and_score(models=models,
                              X_train=X_train,
                              X_test=X_test,
                              y_train=y_train,
                              y_test=y_test)

model_scores

{'KNN': 0.6885245901639344,
 'Logistic_Regression': 0.8688524590163934,
 'Random_Forest': 0.8360655737704918}

```

```

model_compare = pd.DataFrame(model_scores, index=['accuracy']).T.reset_index()
model_compare.columns = ['Model', 'Accuracy']

```

```

fig = px.bar(model_compare,
             x='Model',
             y='Accuracy',
             title='Model Accuracy Comparison',
             text='Accuracy',
             color='Model',
             color_discrete_sequence=px.colors.qualitative.Set2)

fig.update_traces(texttemplate='%{text:.2%}', textposition='outside')
fig.update_layout(yaxis=dict(tickformat=".0%"), showlegend=False)
fig.update_layout(
    yaxis=dict(tickformat=".0%"),
    showlegend=False,
    height=600
)
fig.show()

```



```

train_scores = []
test_scores = []

neighbors = list(range(1, 30))

knn = KNeighborsClassifier()

for i in neighbors:
    knn.set_params(n_neighbors = i)

    knn.fit(X_train, y_train)

    train_scores.append(knn.score(X_train, y_train))

    test_scores.append(knn.score(X_test, y_test))

```

```
train_scores
```

```
[1.0,
 0.8099173553719008,
 0.7727272727272727,
 0.743801652892562,
 0.7603305785123967,
 0.7520661157024794,
 0.743801652892562,
 0.7231404958677686,
 0.71900826446281,
 0.6942148760330579,
 0.7272727272727273,
 0.6983471074380165,
 0.6900826446280992,
 0.6942148760330579,
 0.6859504132231405,
 0.6735537190082644,
 0.6859504132231405,
 0.6652892561983471,
 0.6818181818181818,
 0.6694214876033058,
 0.6859504132231405,
 0.6694214876033058,
 0.7024793388429752,
 0.6735537190082644,
 0.6983471074380165,
 0.6942148760330579,
 0.6983471074380165,
 0.6859504132231405,
 0.6818181818181818]
```

```
test_scores
```

```
[0.6229508196721312,
 0.639344262295082,
 0.6557377049180327,
 0.6721311475409836,
 0.6885245901639344,
 0.7213114754098361,
 0.7049180327868853,
 0.6885245901639344,
 0.6885245901639344,
 0.7049180327868853,
 0.7540983606557377,
 0.7377049180327869,
 0.7377049180327869,
 0.7377049180327869,
 0.6885245901639344,
 0.7213114754098361,
 0.6885245901639344,
 0.6885245901639344,
 0.7049180327868853,
 0.6557377049180327,
 0.7049180327868853,
 0.7213114754098361,
 0.7213114754098361,
 0.7213114754098361,
 0.7213114754098361,
 0.7049180327868853,
 0.7213114754098361,
 0.7213114754098361,
 0.7049180327868853,
 0.7213114754098361]
```

```
fig = go.Figure()
```

```
fig.add_trace(go.Scatter(
    x=neighbors,
    y=train_scores,
    mode='lines+markers',
    name='Train score',
    line=dict(color='blue'),
    marker=dict(size=8)
))
```

```
fig.add_trace(go.Scatter(
    x=neighbors,
    y=test_scores,
    mode='lines+markers',
    name='Test score',
    line=dict(color='orange'),
    marker=dict(size=8)
))
```

```

))

fig.update_layout(
    title="KNN Model Accuracy vs. Number of Neighbors",
    xaxis=dict(title="Number of Neighbors", tickmode='linear', dtick=1),
    yaxis=dict(title="Model Score"),
    width=800,
    height=500,
    legend=dict(title="Legend"),
)


fig.show()

print(f"Maximum KNN score on the test data: {max(test_scores)*100:.2f}%")
knn_accuracy = round(max(test_scores) * 100, 2)

```



KNN Model Accuracy vs. Number of Neighbors



Maximum KNN score on the test data: 75.41%

```

# using randomizedsearchcv for model params tuning
log_reg_grid = {"C": np.logspace(-4, 4, 20),
                "solver": ["liblinear"]}

rf_grid = {"n_estimators": np.arange(10, 1000, 50),
           "max_depth": [None, 3, 5, 10],
           "min_samples_split": np.arange(2, 20, 2),
           "min_samples_leaf": np.arange(1, 20, 2)}

```

```

#logistic regression
np.random.seed(42)

rs_log_reg = RandomizedSearchCV(LogisticRegression(),
                                param_distributions=log_reg_grid,
                                cv=5,
                                n_iter=20,
                                verbose=True)

rs_log_reg.fit(X_train, y_train);

```



Fitting 5 folds for each of 20 candidates, totalling 100 fits

```
rs_log_reg.best_params_
```



```
{'solver': 'liblinear', 'C': np.float64(0.23357214690901212)}
```

```
log_reg_accuracy = round(rs_log_reg.score(X_test, y_test)*100,2)
```

```

#random forest
np.random.seed(42)

rs_rf = RandomizedSearchCV(RandomForestClassifier(),
                            param_distributions=rf_grid,
                            cv=5,


```

```

        n_iter=20,
        verbose=True)

```

```
rs_rf.fit(X_train, y_train);
```


 Fitting 5 folds for each of 20 candidates, totalling 100 fits

```

# Find the best parameters
rs_rf.best_params_

```

```

 {'n_estimators': np.int64(210),
  'min_samples_split': np.int64(4),
  'min_samples_leaf': np.int64(19),
  'max_depth': 3}

```

```

# Evaluating the randomized search random forest model
rf_accuracy = round(rs_rf.score(X_test, y_test)*100,2)

```

```

model_scores = {
    "KNN": knn_accuracy,
    "logistic regression":log_reg_accuracy,
    "random forest":rf_accuracy
}

```

```

model_compare = pd.DataFrame(model_scores, index=['accuracy']).T.reset_index()
model_compare.columns = ['Model', 'Accuracy']

```

```

fig = px.bar(model_compare,
             x='Model',
             y='Accuracy',
             title='Model Accuracy Comparison after tuning',
             text='Accuracy',
             color='Model',
             color_discrete_sequence=px.colors.qualitative.Set2)

```


```

fig.update_layout(
    yaxis=dict(tickformat=".-0%"),
    showlegend=False,
    height=600
)
fig.show()

```



Model Accuracy Comparison after tuning



```

y_preds = rs_log_reg.predict(X_test)
y_preds


```

```
array([[0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0,
        0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0]])
```

y\_test

```
array([[0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0,
        0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0]])
```

```
#model evaluation
RocCurveDisplay.from_estimator(rs_log_reg, X_test, y_test)
plt.show()
```



```
print(confusion_matrix(y_test, y_preds))
```

```
[[25  4]
 [ 3 29]]
```


```
sns.set(font_scale=1.5)
```

```
def plot_conf_mat(y_test, y_preds):
```

```
    fig, ax = plt.subplots(figsize=(3, 3))
    ax = sns.heatmap(confusion_matrix(y_test, y_preds),
                     annot=True,
                     cbar=False)
    plt.xlabel("True label")
    plt.ylabel("Predicted label")

    bottom, top = ax.get_ylim()
    ax.set_ylim(bottom, top)
    plt.show()
```

```
plot_conf_mat(y_test, y_preds)
```



```
print(classification_report(y_test, y_preds))
```

```
↵
```

	precision	recall	f1-score	support
0	0.89	0.86	0.88	29
1	0.88	0.91	0.89	32
accuracy			0.89	61
macro avg	0.89	0.88	0.88	61
weighted avg	0.89	0.89	0.89	61

```
#instanciating log_reg wit best params to check which feature helps most in model training
```

```
clf = LogisticRegression(C=0.23357214690901212,  
                        solver="liblinear")
```